

Homework 6

PSTAT 131/231

Tree-Based Models

```
library(tidyverse)
library(tidymodels)
library(rpart.plot)
library(vip)
library(janitor)
library(randomForest)
library(xgboost)
library(ranger)
tidymodels_prefer()
```

Exercise 1

Read in the data and set things up as in Homework 5:

- Use `clean_names()`
- Filter out the rarer Pokémon types
- Convert `type_1` and `legendary` to factors

Do an initial split of the data; you can choose the percentage for splitting. Stratify on the outcome variable.

Fold the training set using v -fold cross-validation, with $v = 5$. Stratify on the outcome variable.

Set up a recipe to predict `type_1` with `legendary`, `generation`, `sp_atk`, `attack`, `speed`, `defense`, `hp`, and `sp_def`:

- Dummy-code `legendary` and `generation`;
- Center and scale all predictors.

```
pokemon <- read.csv("Pokemon.csv")
pokemon_clean <- clean_names(pokemon)

pokemon_filter <- pokemon_clean %>%
  filter(type_1 == c("Bug", "Fire", "Grass", "Normal", "Water", "Psychic")) %>%
  mutate_at(vars(type_1, legendary), factor)

set.seed(2022)

pokemon_split <- initial_split(pokemon_filter, prop = 0.70,
                              strata = type_1)
pokemon_train <- training(pokemon_split)
pokemon_test <- testing(pokemon_split)

dim(pokemon_train)
```

```
## [1] 54 13
dim(pokemon_test)

## [1] 28 13
pokemon_folds <- vfold_cv(pokemon_train, v = 5, strata = type_1)
pokemon_folds

## # 5-fold cross-validation using stratification
## # A tibble: 5 x 2
##   splits      id
##   <list>      <chr>
## 1 <split [40/14]> Fold1
## 2 <split [42/12]> Fold2
## 3 <split [43/11]> Fold3
## 4 <split [45/9]>  Fold4
## 5 <split [46/8]>  Fold5

pokemon_recipe <- recipe(type_1 ~ legendary + generation + sp_atk + attack
  + speed + defense + hp + sp_def, data = pokemon_train) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_center(all_numeric_predictors()) %>%
  step_scale(all_numeric_predictors())

pokemon_recipe

## Recipe
##
## Inputs:
##
##   role #variables
##   outcome      1
##   predictor      8
##
## Operations:
##
## Dummy variables from all_nominal_predictors()
## Centering for all_numeric_predictors()
## Scaling for all_numeric_predictors()
```

Exercise 2

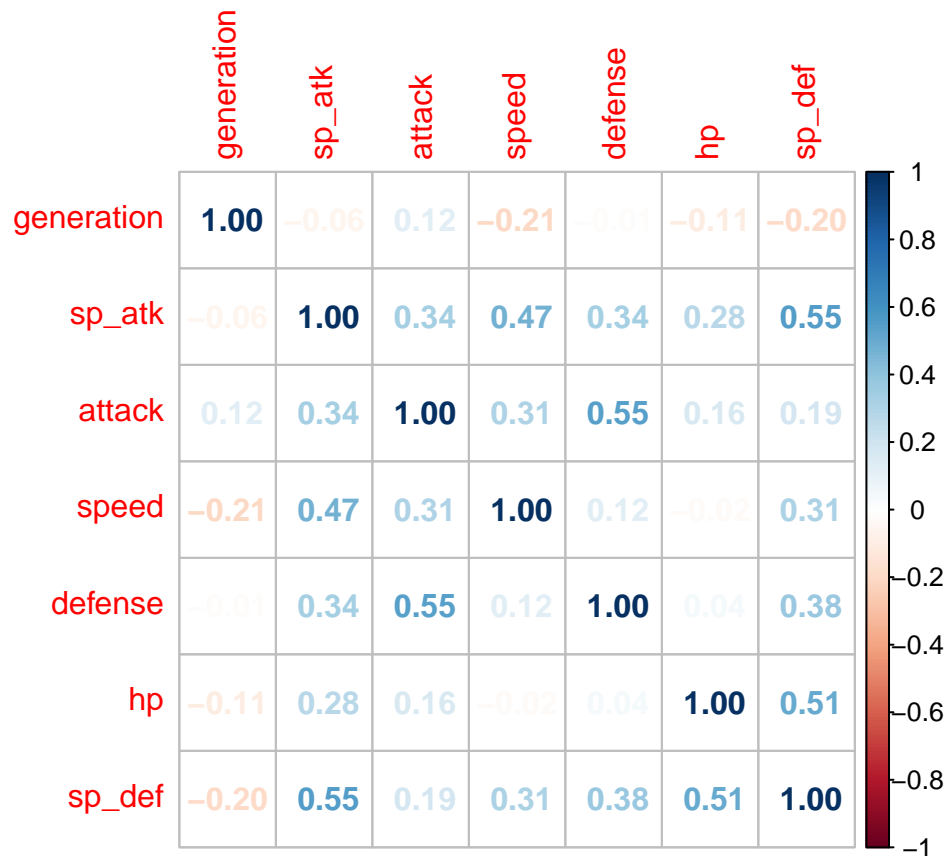
Create a correlation matrix of the training set, using the `corrplot` package. *Note: You can choose how to handle the continuous variables for this plot; justify your decision(s).*

What relationships, if any, do you notice? Do these relationships make sense to you?

```
library(corrplot)

cor_pokemon <- pokemon_train %>%
  select(c(generation, sp_atk, attack, speed, defense, hp, sp_def)) %>%
  cor()

corrplot(cor_pokemon, method = "number")
```



Exercise 3

First, set up a decision tree model and workflow. Tune the `cost_complexity` hyperparameter. Use the same levels we used in Lab 7 – that is, `range = c(-3, -1)`. Specify that the metric we want to optimize is `roc_auc`.

Print an `autoplot()` of the results. What do you observe? Does a single decision tree perform better with a smaller or larger complexity penalty?

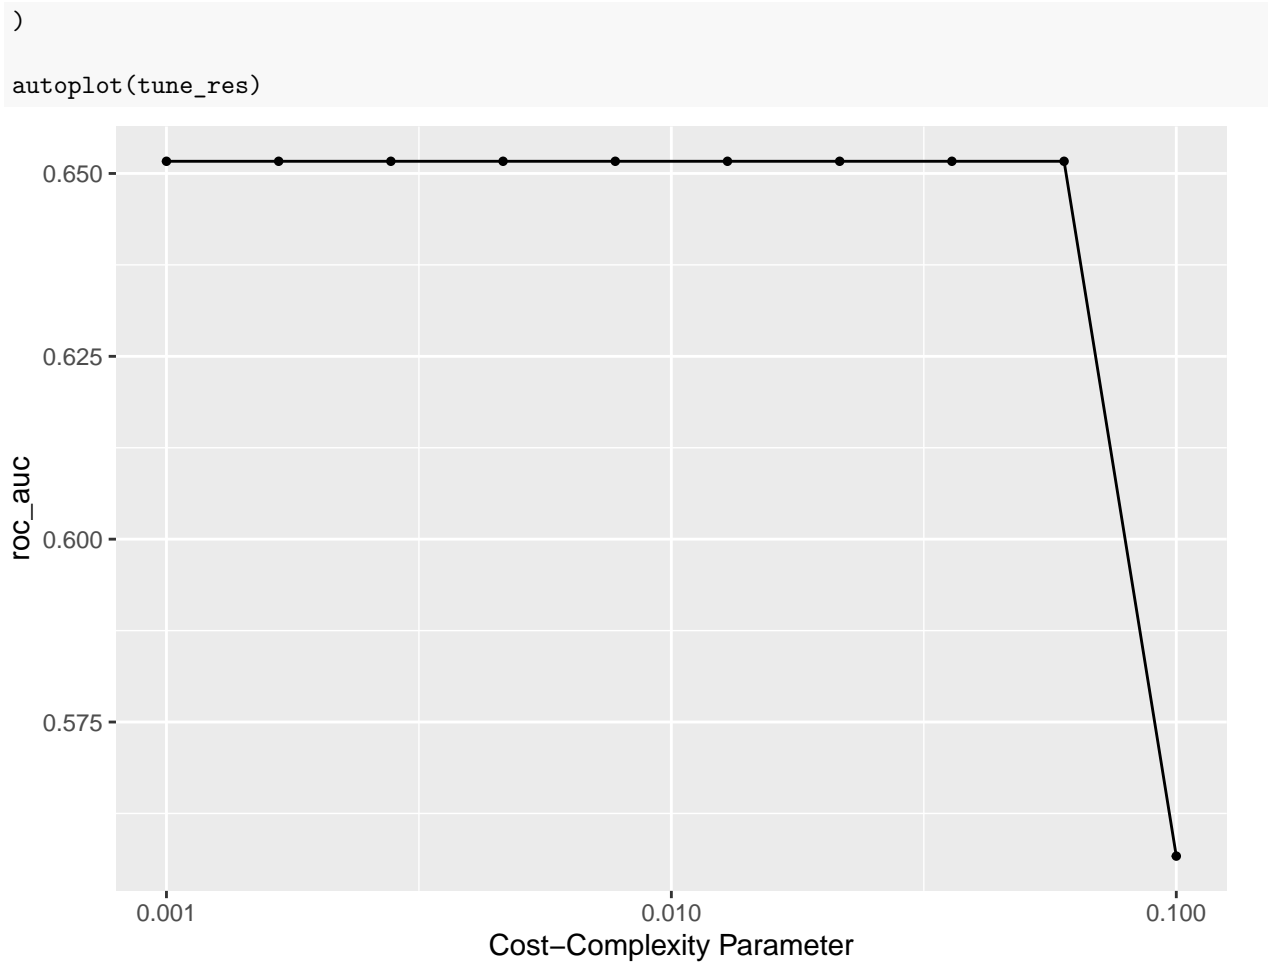
```
tree_spec <- decision_tree() %>%
  set_engine("rpart")

class_tree_spec <- tree_spec %>%
  set_mode("classification")

class_tree_wf <- workflow() %>%
  add_model(class_tree_spec %>% set_args(cost_complexity = tune())) %>%
  add_recipe(pokemon_recipe)

param_grid <- grid_regular(cost_complexity(range = c(-3, -1)), levels = 10)

tune_res <- tune_grid(
  class_tree_wf,
  resamples = pokemon_folds,
  grid = param_grid,
  metrics = metric_set(roc_auc)
```



Solution: A single decision tree performs better with a smaller complexity penalty.

Exercise 4

What is the `roc_auc` of your best-performing pruned decision tree on the folds? *Hint: Use `collect_metrics()` and `arrange()`.*

```
collect_metrics(tune_res) %>%
  arrange(-mean)
```

```
## # A tibble: 10 x 7
```

	cost_complexity	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
## 1	0.001	roc_auc	hand_till	0.652	5	0.0465	Preprocessor1_Model01
## 2	0.00167	roc_auc	hand_till	0.652	5	0.0465	Preprocessor1_Model02
## 3	0.00278	roc_auc	hand_till	0.652	5	0.0465	Preprocessor1_Model03
## 4	0.00464	roc_auc	hand_till	0.652	5	0.0465	Preprocessor1_Model04
## 5	0.00774	roc_auc	hand_till	0.652	5	0.0465	Preprocessor1_Model05
## 6	0.0129	roc_auc	hand_till	0.652	5	0.0465	Preprocessor1_Model06
## 7	0.0215	roc_auc	hand_till	0.652	5	0.0465	Preprocessor1_Model07
## 8	0.0359	roc_auc	hand_till	0.652	5	0.0465	Preprocessor1_Model08
## 9	0.0599	roc_auc	hand_till	0.652	5	0.0465	Preprocessor1_Model09
## 10	0.1	roc_auc	hand_till	0.557	5	0.0314	Preprocessor1_Model10

Solution: The roc_auc of best-performing pruned decision tree on the folds is 0.651667.

Exercise 5

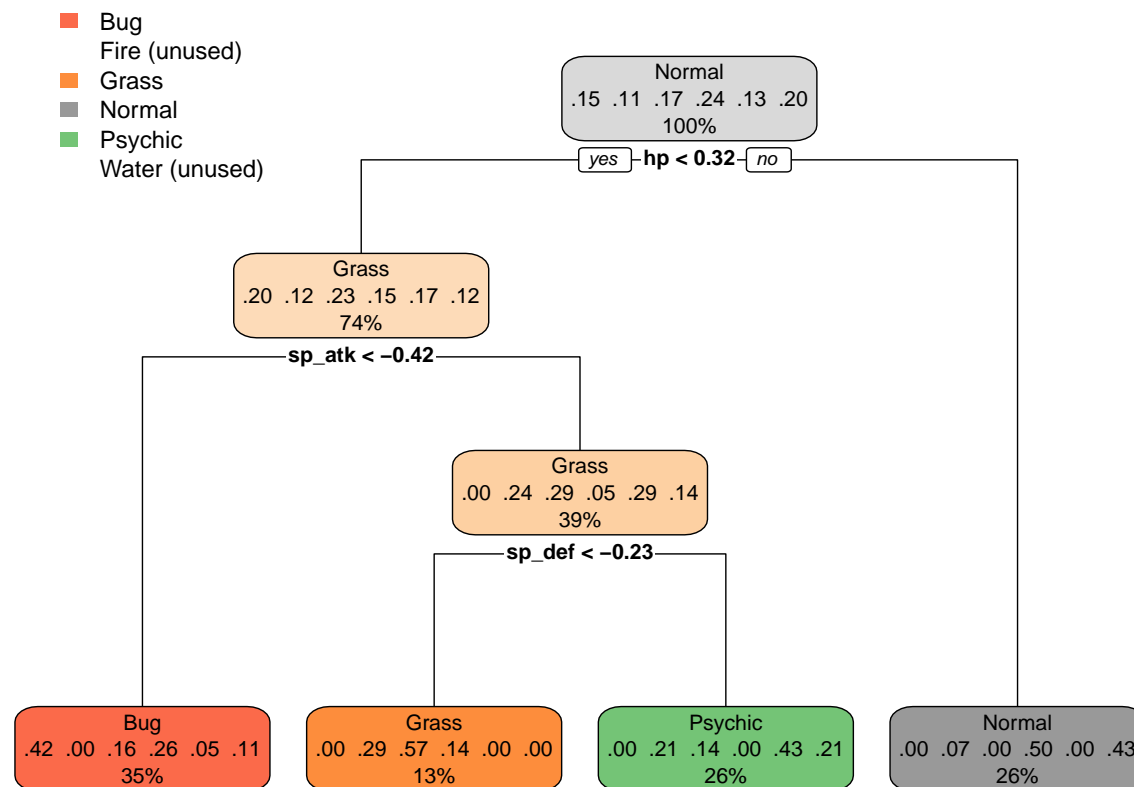
Using `rpart.plot`, fit and visualize your best-performing pruned decision tree with the *training* set.

```
best_complexity <- select_best(tune_res, metric = "roc_auc")

class_tree_final <- finalize_workflow(class_tree_wf, best_complexity)

class_tree_final_fit <- fit(class_tree_final, data = pokemon_train)

class_tree_final_fit %>%
  extract_fit_engine() %>%
  rpart.plot()
```



Exercise 5

Now set up a random forest model and workflow. Use the `ranger` engine and set `importance = "impurity"`. Tune `mtry`, `trees`, and `min_n`. Using the documentation for `rand_forest()`, explain in your own words what each of these hyperparameters represent.

Create a regular grid with 8 levels each. You can choose plausible ranges for each hyperparameter. Note that `mtry` should not be smaller than 1 or larger than 8. Explain why not. What type of model would `mtry = 8` represent?

```
rf_spec <- rand_forest() %>%
  set_engine("ranger", importance = "impurity") %>%
```

```

set_mode("classification")

rf_wf <- workflow() %>%
  add_model(rf_spec %>% set_args(mtry = tune(), trees = tune(), min_n = tune())) %>%
  add_recipe(pokemon_recipe)

rf_grid <- grid_regular(mtry(range = c(1, 8)), trees(range = c(1, 8)),
  min_n(range = c(1, 8)), levels = 8)

```

Solution: mtry: the number of predictors that will be randomly sampled.

trees: the number of trees contained in the ensemble.

min_n: the minimum number of data points in a node that are required for the node to be split further.

Because there are only 8 predictors thus $1 \leq mtry \leq 8$. mtry = 8 means we use all the predictors to be randomly sampled at each split when creating the tree models.

Exercise 6

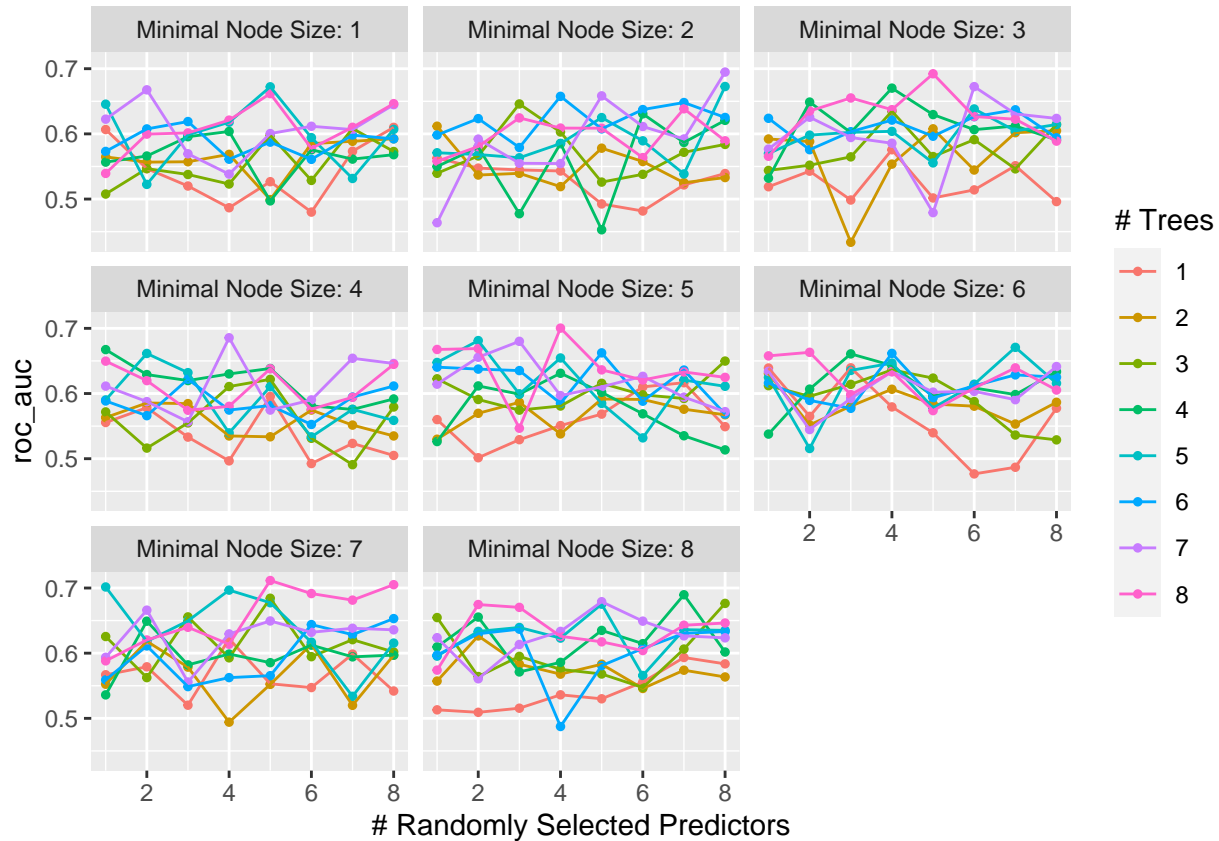
Specify roc_auc as a metric. Tune the model and print an autoplot() of the results. What do you observe? What values of the hyperparameters seem to yield the best performance?

```

rf_tune_res <- tune_grid(
  rf_wf,
  resamples = pokemon_folds,
  grid = rf_grid,
  metrics = metric_set(roc_auc)
)

autoplot(rf_tune_res)

```



Solution: The relationship is complex. trees = 8, mtyes = 5, min_n = 7 seems to yield the best performance.

Exercise 7

What is the roc_auc of your best-performing random forest model on the folds? *Hint: Use collect_metrics() and arrange().*

```
collect_metrics(rf_tune_res) %>%
  arrange(-mean)
```

```
## # A tibble: 512 x 9
##   mtry trees min_n .metric .estimator mean    n std_err .config
##   <int> <int> <int> <chr>   <chr>   <dbl> <int>   <dbl> <chr>
## 1     5     8     7   roc_auc hand_till 0.711     5 0.0230 Preprocessor1_Model~
## 2     8     8     7   roc_auc hand_till 0.705     5 0.0407 Preprocessor1_Model~
## 3     1     5     7   roc_auc hand_till 0.702     5 0.0292 Preprocessor1_Model~
## 4     4     8     5   roc_auc hand_till 0.700     5 0.0426 Preprocessor1_Model~
## 5     4     5     7   roc_auc hand_till 0.697     5 0.0271 Preprocessor1_Model~
## 6     8     7     2   roc_auc hand_till 0.695     5 0.0350 Preprocessor1_Model~
## 7     5     8     3   roc_auc hand_till 0.692     5 0.0118 Preprocessor1_Model~
## 8     6     8     7   roc_auc hand_till 0.691     5 0.0262 Preprocessor1_Model~
## 9     7     4     8   roc_auc hand_till 0.690     5 0.0156 Preprocessor1_Model~
## 10    4     7     4   roc_auc hand_till 0.686     5 0.0420 Preprocessor1_Model~
## # ... with 502 more rows
```

Solution: The roc_auc of your best-performing random forest model on the folds is 0.7114815.

Exercise 8


Create a variable importance plot, using `vip()`, with your best-performing random forest model fit on the *training* set.

```
best_complexity <- select_best(rf_tune_res, metric = "roc_auc")

rf_final <- finalize_workflow(rf_wf, best_complexity)

rf_final_fit <- fit(rf_final, data = pokemon_train)

rf_final_fit %>%
  pull_workflow_fit() %>%
  vip()
```



Which variables were most useful? Which were least useful? Are these results what you expected, or not?

Solution: `sp_atk` is most useful and `legendary` is least useful. Yes, the `sp_atk` is the most important character of the type of pokemons.


Exercise 9

Finally, set up a boosted tree model and workflow. Use the `xgboost` engine. Tune `trees`. Create a regular grid with 10 levels; let `trees` range from 10 to 2000. Specify `roc_auc` and again print an `autoplot()` of the results.

What do you observe?

What is the roc_auc of your best-performing boosted tree model on the folds? *Hint: Use collect_metrics() and arrange().*

```
boost_spec <- boost_tree() %>%  
  set_engine("xgboost") %>%  
  set_mode("classification")  
  
boost_wf <- workflow() %>%  
  add_model(boost_spec %>% set_args(trees = tune())) %>%  
  add_recipe(pokemon_recipe)  
  
boost_grid <- grid_regular(trees(range = c(10, 2000)), levels = 10)  
  
boost_tune_res <- tune_grid(  
  boost_wf,  
  resamples = pokemon_folds,  
  grid = boost_grid,  
  metrics = metric_set(roc_auc)  
)  
  
autoplot(boost_tune_res)
```



```
collect_metrics(boost_tune_res) %>%  
  arrange(desc(mean), by_group = TRUE)
```

```
## # A tibble: 10 x 7  
##   trees .metric .estimator mean      n std_err .config  
##   <int> <chr>   <chr>      <dbl> <int>  <dbl> <chr>
```

```
## 1 1115 roc_auc hand_till 0.649 5 0.0340 Preprocessor1_Model06
## 2 673 roc_auc hand_till 0.648 5 0.0336 Preprocessor1_Model04
## 3 1336 roc_auc hand_till 0.648 5 0.0336 Preprocessor1_Model07
## 4 1557 roc_auc hand_till 0.646 5 0.0337 Preprocessor1_Model08
## 5 1778 roc_auc hand_till 0.646 5 0.0337 Preprocessor1_Model09
## 6 2000 roc_auc hand_till 0.646 5 0.0337 Preprocessor1_Model10
## 7 452 roc_auc hand_till 0.644 5 0.0344 Preprocessor1_Model03
## 8 894 roc_auc hand_till 0.644 5 0.0338 Preprocessor1_Model05
## 9 231 roc_auc hand_till 0.628 5 0.0397 Preprocessor1_Model02
## 10 10 roc_auc hand_till 0.583 5 0.0370 Preprocessor1_Model01
```

Solution: We observe that with the increase of the trees, the roc_auc increases quickly and slow down and last finally. The roc_auc of the best-performing boosted tree model on the folds is 0.6492593.

Exercise 10

Display a table of the three ROC AUC values for your best-performing pruned tree, random forest, and boosted tree models. Which performed best on the folds? Select the best of the three and use `select_best()`, `finalize_workflow()`, and `fit()` to fit it to the *testing* set.

Print the AUC value of your best-performing model on the testing set. Print the ROC curves. Finally, create and visualize a confusion matrix heat map.

Which classes was your model most accurate at predicting? Which was it worst at?

```
roc_auc_pre <- c(0.6516667, 0.7114815, 0.6492593)
models <- c("best-performing pruned tree model", "random forest model", "boosted tree model")
results <- tibble(roc_auc_pre = roc_auc_pre, models = models)

results %>%
  arrange(-roc_auc_pre)

## # A tibble: 3 x 2
##   roc_auc_pre models
##   <dbl> <chr>
## 1 0.711 random forest model
## 2 0.652 best-performing pruned tree model
## 3 0.649 boosted tree model

best_complexity <- select_best(rf_tune_res, metric = "roc_auc")

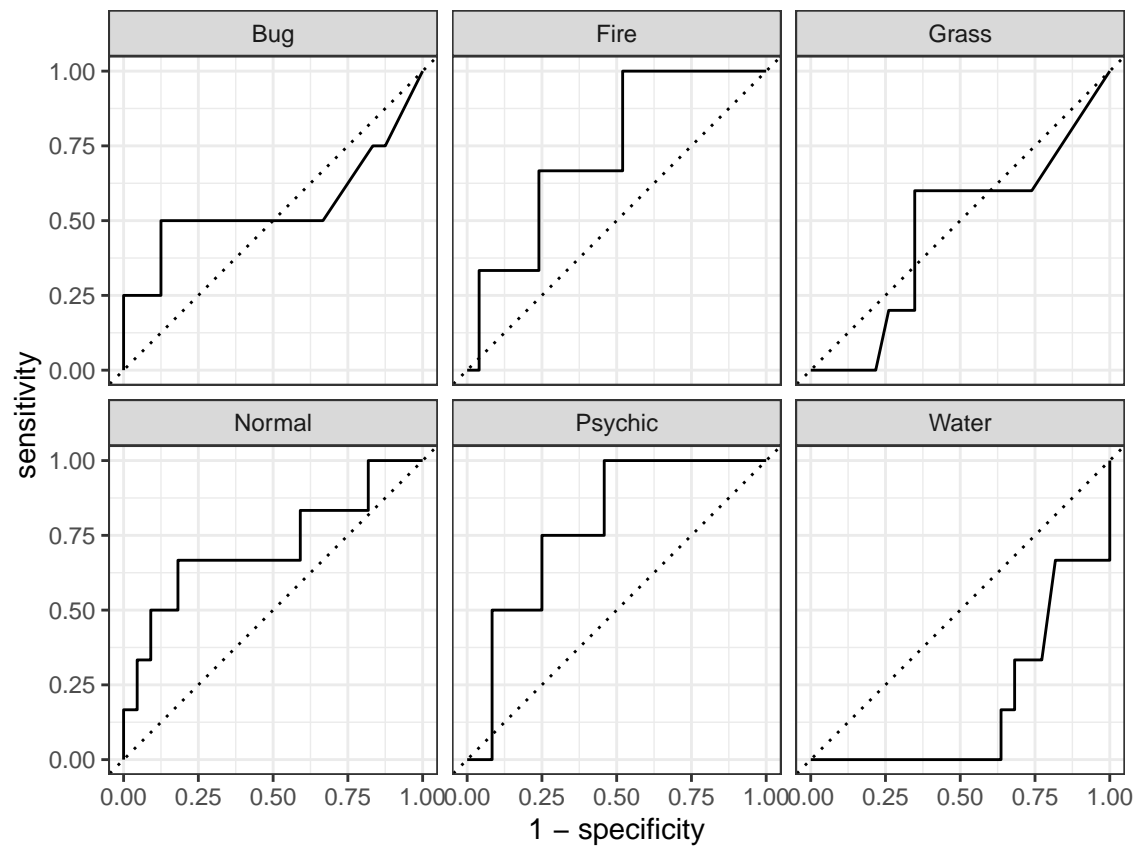
rf_final <- finalize_workflow(rf_wf, best_complexity)

rf_final_fit <- fit(rf_final, data = pokemon_train)

augment(rf_final_fit, new_data = pokemon_test) %>%
  roc_auc(type_1, .pred_Bug:.pred_Water)

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr> <chr> <dbl>
## 1 roc_auc hand_till 0.573

augment(rf_final_fit, new_data = pokemon_test) %>%
  roc_curve(type_1, .pred_Bug:.pred_Water) %>%
  autoplot()
```



```
augment(rf_final_fit , new_data = pokemon_test) %>%
  conf_mat(type_1, .pred_class) %>%
  autoplot(type = "heatmap")
```

Prediction	Bug -	1	0	0	0	0	2
	Fire -	1	1	1	0	1	0
	Grass -	0	2	0	0	1	0
	Normal -	1	0	1	4	0	3
	Psychic -	1	0	1	1	1	1
	Water -	0	0	2	1	1	0
		Bug	Fire	Grass	Normal	Psychic	Water
		Truth					

Solution: The random forest model performed best on the folds. The model is best at predicting fire and is worst at predicting grass.